For best experience please consider to run **Hosam_Mahmoud_300327269_Reinforcement_Learning.py** on Spyder to see all created variables on variables explorer tab.

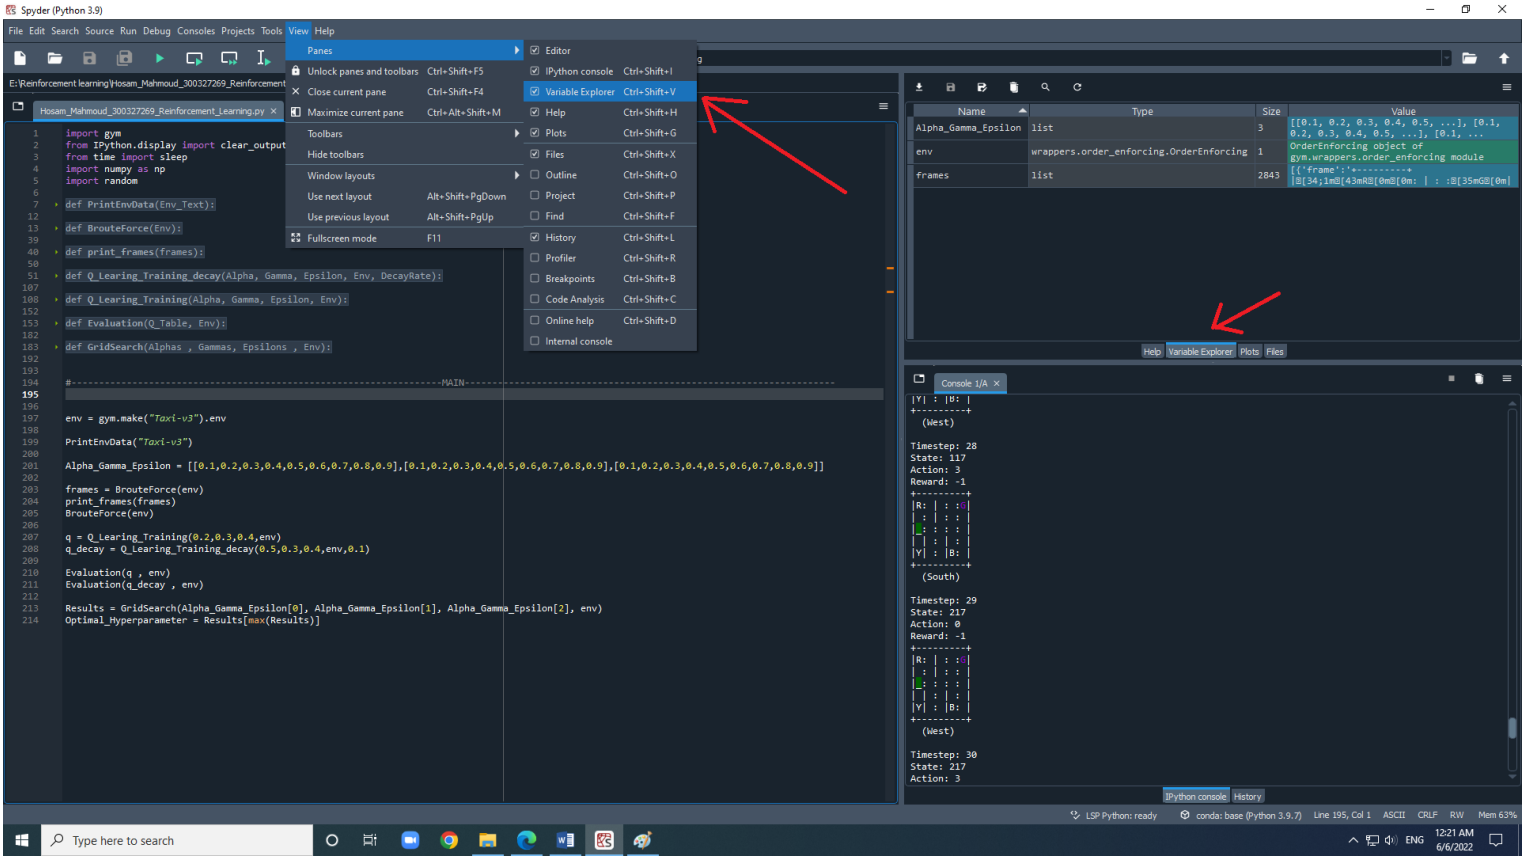

Please make sure that you have installed these libraries before importing it ->

import gym
from IPython.display import clear_output
from time import sleep
import numpy as np
import random

To install the missing libraries, you have to write something like this ->

>>> pip install *numpy*

In Spyder console.

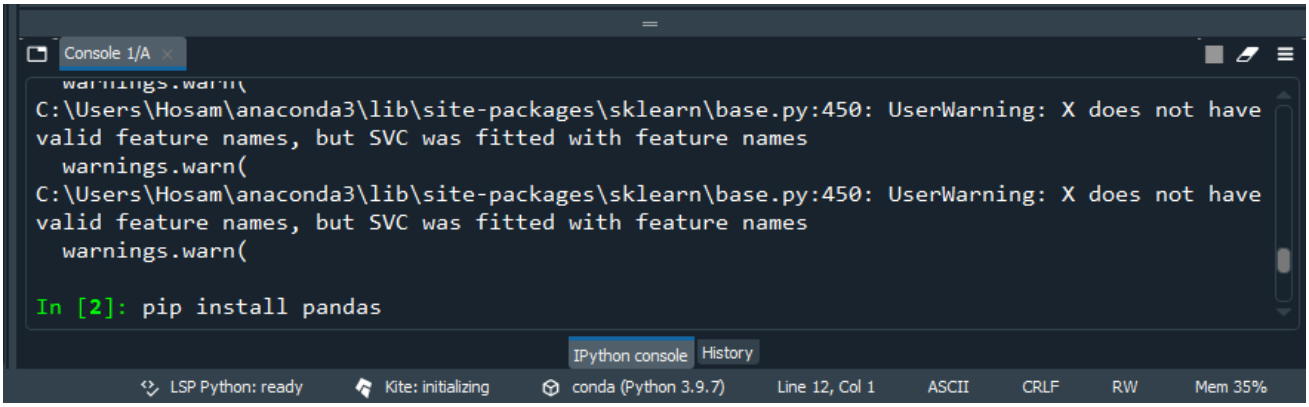

or

>>> conda install pandas

In Anaconda prompt.

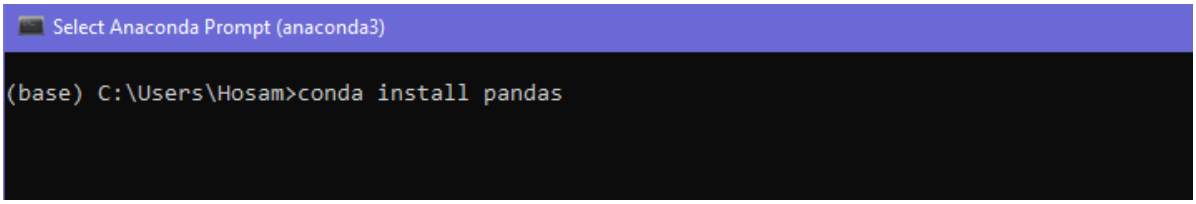

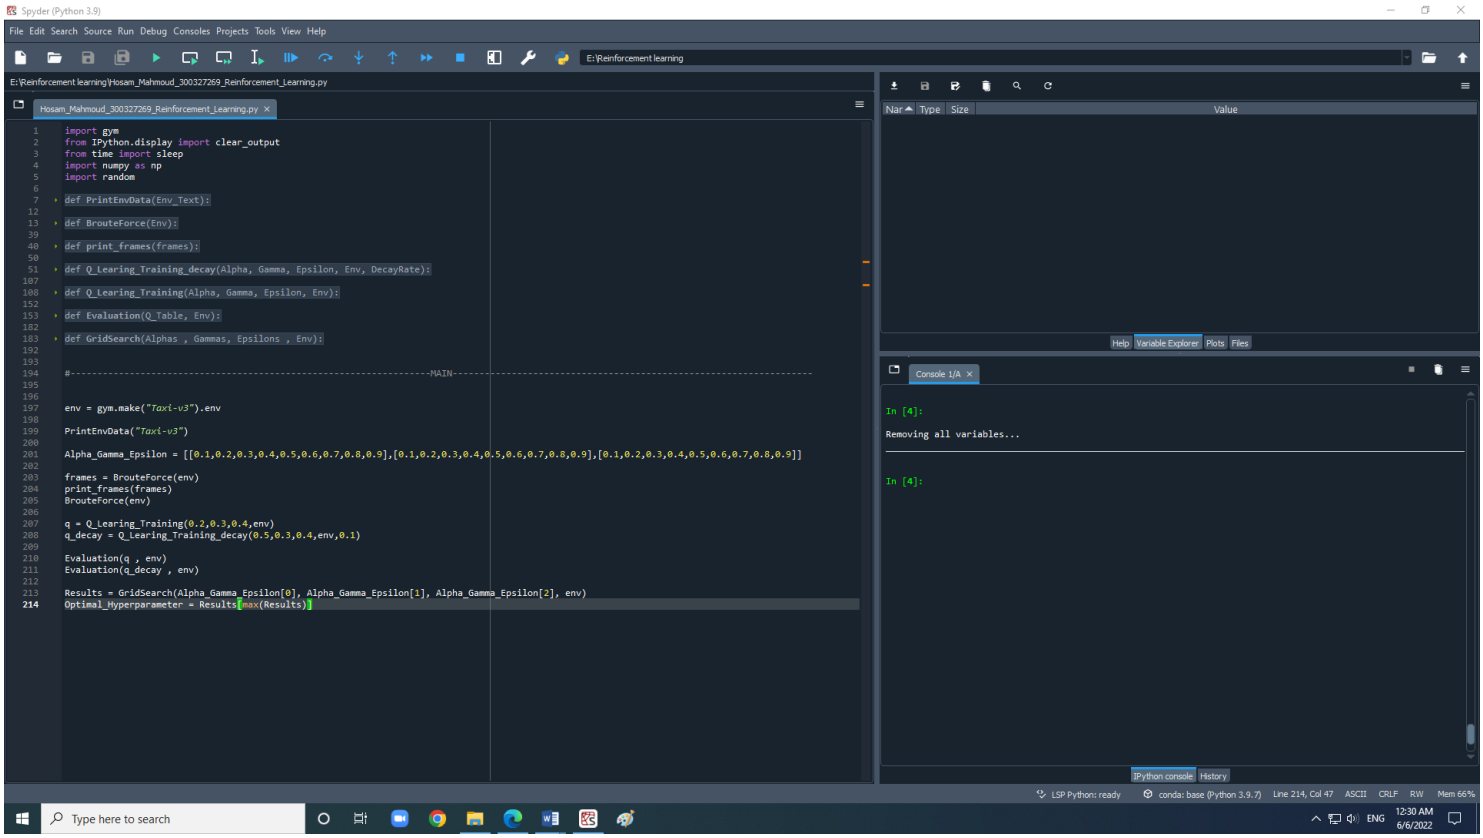

Here I have 7 functions in my code.

def PrintEnvData(Env_Text): is for printing the number of actions and states of the environments.

def BrouteForce(Env): to try all the possible paths for one state.

def print_frames(frames): to animate the paths.

def Q_Learing_Training_decay(Alpha, Gamma, Epsilon, Env, DecayRate): Training function using a decay over episodes.

Here is the parameter tuning part, but in case of **gamma** and **epsilon** I have tried to change its values but I got an infinite loop when applying evaluation.

```
        if i % 100 == 0:
            clear_output(wait = True)
            print(f"Episode: {i}")
            alpha = abs(alpha - (1/(1 + (DecayRate * 100000))) * alpha)
            # gamma = abs(gamma - (1/(1 + (DecayRate * 100000))) * gamma)
            # epsilon = abs(epsilon - (1/(1 + (DecayRate * 100000))) * epsilon)

            print('Alpha = ', alpha)
            print('Gamma = ', gamma)
            print('Epsilon = ', epsilon)

            alpha = Alpha if alpha == 0 else alpha
            # gamma = Gamma if gamma == 0 else gamma
            # epsilon = Epsilon if epsilon == 0 else epsilon

    print("Training finished.\n")
    return q_table
```

def Q_Learing_Training(Alpha, Gamma, Epsilon, Env): Training function without parameter tuning.
def Evaluation(Q_Table, Env): to return the performance indicator of each reinforcement model.
def GridSearch(Alphas , Gammas, Epsilons , Env): to apply GridSearch and optain the optimal hyperparameters.

```
193
194  #-----------------------------------------------------------------MAIN-----------------------------------------------------------------
195
196
197  env = gym.make("Taxi-v3").env
198
199  PrintEnvData("Taxi-v3")
200
201  Alpha_Gamma_Epsilon = [[0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9],[0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9],[0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9]]
202
203  frames = BrouteForce(env)
204  print_frames(frames)
205  BrouteForce(env)
206
207  q = Q_Learing_Training(0.2,0.3,0.4,env)
208  q_decay = Q_Learing_Training_decay(0.5,0.3,0.4,env,0.1)
209
210  Evaluation(q , env)
211  Evaluation(q_decay , env)
212
213  Results = GridSearch(Alpha_Gamma_Epsilon[0], Alpha_Gamma_Epsilon[1], Alpha_Gamma_Epsilon[2], env)
214  Optimal_Hyperparameter = Results[max(Results)]
```

For the MAIN part I have created environment and print its number of states and actions.

```
In [7]: runfile('E:/Reinforcement learning/Hosam_Mahmoud_300327269_Reinforcement_Learning.py', wdir='E:/
Reinforcement learning')
Action Space Discrete(6)
State Space Discrete(500)

In [8]:
```

After that I have created list of each parameter (alpha, gamma and epsilon), and I'll use it for grid search **->** 9*9*9 but this toke too long time, so I tried with another lists **->** [0.3,0.6,0.9] which will give us a total number of combinations = 3*3*3

And after that I animated the paths of Bruteforce function.

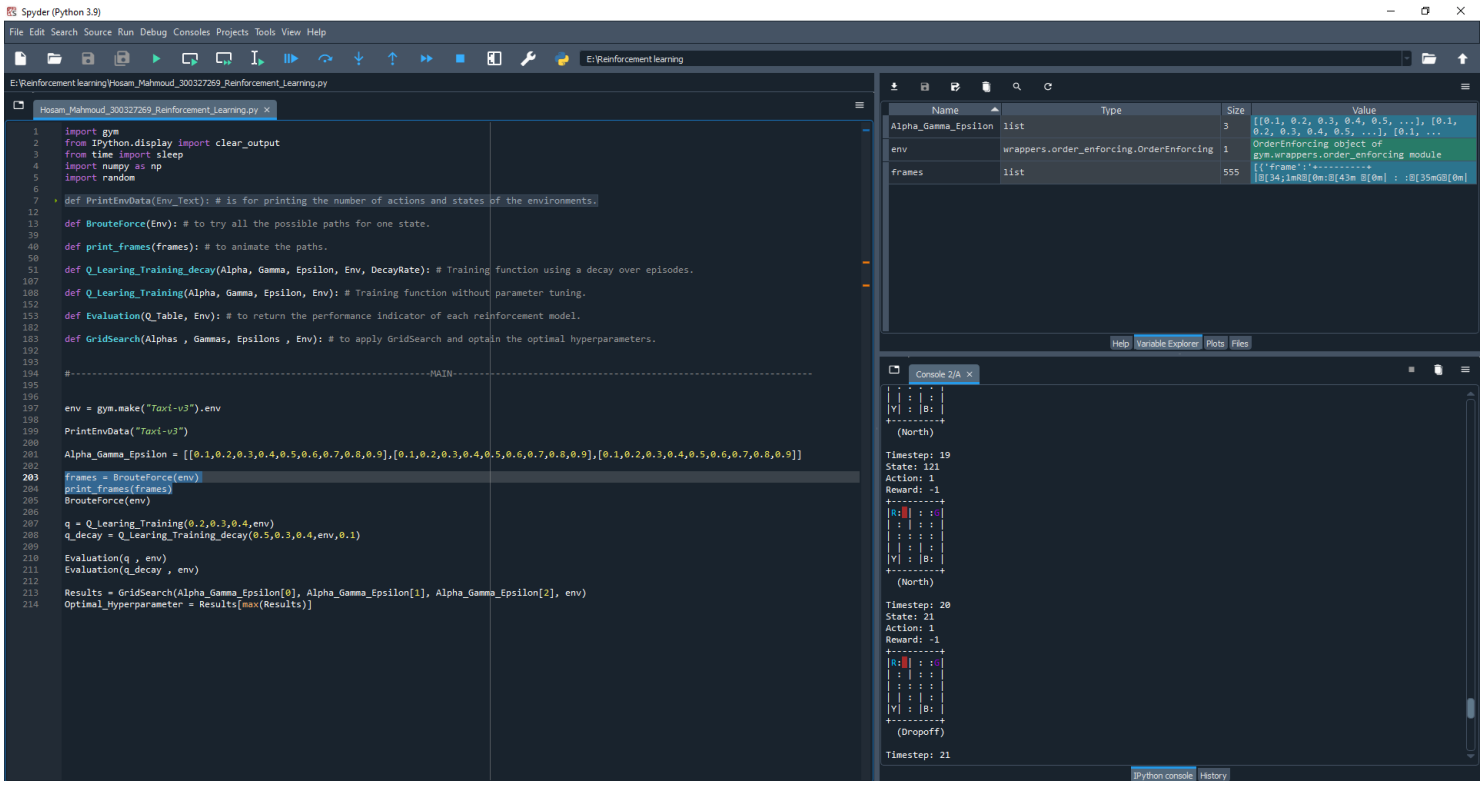

After that I trained the two models (using decay and without using decay).
And the values of alpha are changed every 100 episodes.

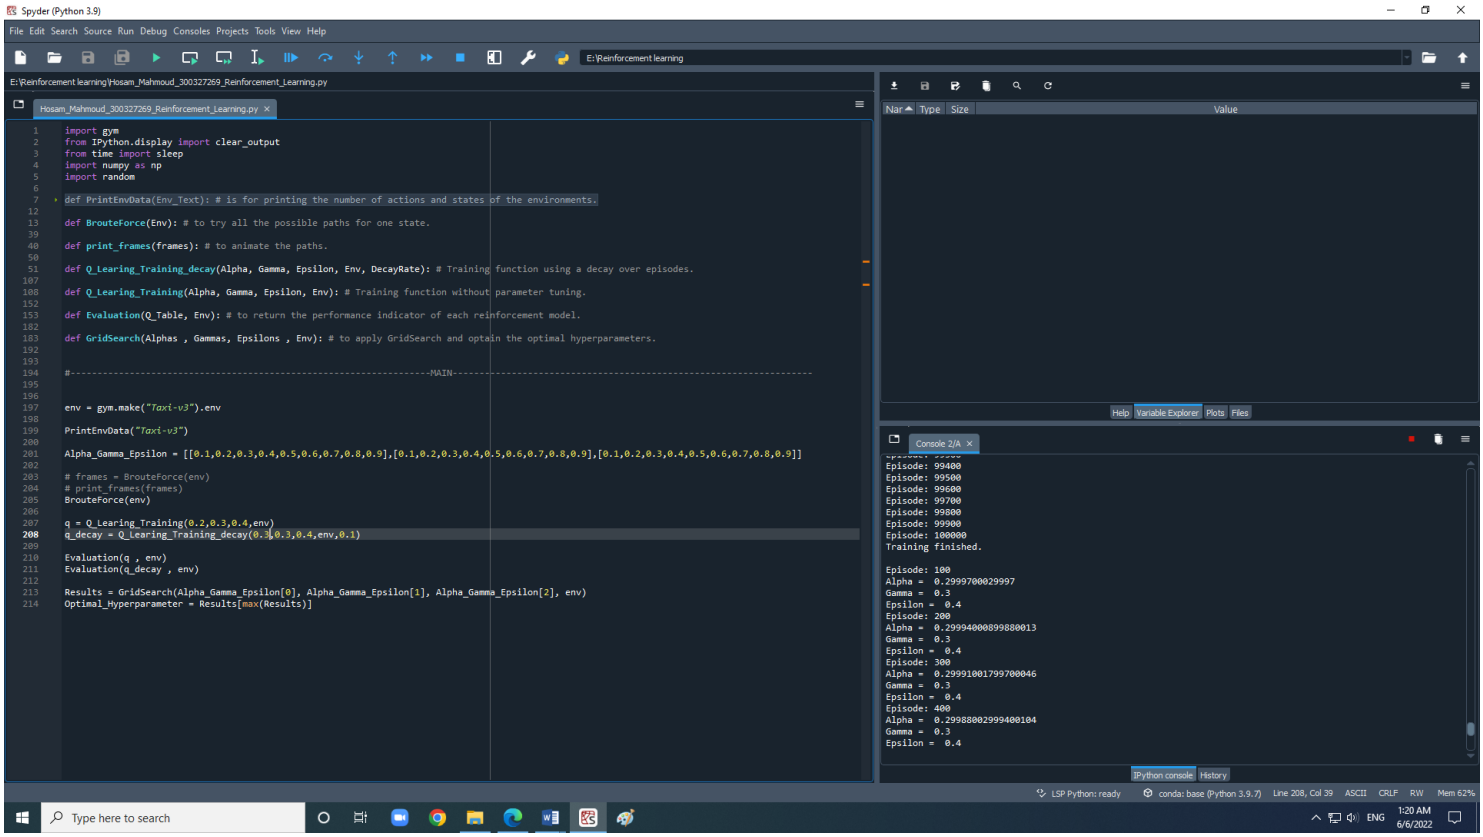

After Evaluation (without using decay)

```
Results after 1000 episodes:
Average timesteps per episode: 13.284
Average penalties per episode: 0.0
Average reward per episode: 20.0
```

After Evaluation (using decay)

```
Results after 1000 episodes:
Average timesteps per episode: 12.976
Average penalties per episode: 0.0
Average reward per episode: 20.0
```

And we found that (with using decay) is faster than the normal model.

After that I applied GridSearch on function (without using Decay)

```
182
183     def GridSearch(Alphas , Gammas, Epsilons , Env): # to apply GridSearch and optain the optimal hyperparameters.
184         Dictionary = {}
185         for Alpha in Alphas:
186             for Gamma in Gammas:
187                 for Epsilon in Epsilons:
188                     q = Q_Learing_Training(Alpha, Gamma , Epsilon, Env)
189                     Metric = Evaluation(q , Env)
190                     Dictionary[Metric] = [Alpha, Gamma , Epsilon]
191         return Dictionary
```

For experiments purposes I have choose small values for parameters so the program run faster -> 2*2*2

Alpha_Gamma_Epsilon - List (3 elements)

Ind	Type	Size	Value
0	list	2	[0.3, 0.6]
1	list	2	[0.3, 0.6]
2	list	2	[0.3, 0.6]

After applying GridSearch we will choose the maximum key and return its parameter And the key is equal to that equation -> **Metric = (total_reward/(total_penalties + total_epochs)).**

Results - Dictionary (7 elements)

Key	Type	Size	Value
1.54762825969202	list	3	[0.3, 0.6, 0.6]
1.54499806875241	list	3	[0.3, 0.3, 0.3]
1.54095076662301	list	3	[0.3, 0.6, 0.3]
1.53586238673015	list	3	[0.6, 0.3, 0.3]
1.52963671128107	list	3	[0.6, 0.6, 0.6]
1.52230172020094	list	3	[0.6, 0.3, 0.6]
1.52021891152326	list	3	[0.6, 0.6, 0.3]

And Finally, we obtained the optimal hyperparameters.

```
Results = GridSearch(Alpha_Gamma_Epsilon[0], Alpha_Gamma_Epsilon[1], Alpha_Gamma_Epsilon[2], env)
Optimal_Hyperparameter = Results[max(Results)]
```

Optimal_Hyperparameter - List (3 elements)

Ind	Type	Size	Value
0	float	1	0.3
1	float	1	0.6
2	float	1	0.6